

Data Mining: Artificial Neural Network


Lecture Notes for Chapter 3 Data Mining

<https://ml-graph.github.io/winter-2025/>


Yu Wang, Ph.D.
yuwang@uoregon.edu
Assistant Professor
Computer Science
University of Oregon
CS 453/553 – Winter 2025

**Course Lecture is very heavily based on
“Introduction to Data Mining”
by Tan, Steinbach, Karpatne, Kumar**


Limitations of Decision-Tree and many other model



Real-world Intuition




Basic Architecture of Perceptron




Artificial Neural Networks (ANN)

- Basic Idea: A complex non-linear function can be learned as a composition of simple processing units
- ANN is a collection of simple processing units (nodes) that are connected by directed links (edges)
 - Every node receives signals from incoming edges, performs computations, and transmits signals to outgoing edges
 - Analogous to *human brain* where nodes are neurons and signals are electrical impulses
 - Weight of an edge determines the strength of connection between the nodes
- Simplest ANN: Perceptron (single neuron)

Basic Architecture of Perceptron



$$y = \sigma(w^T x + b)$$



**What happens if
there is no nonlinear
activation?**

Linear Regression

- Data - $\{(x^{(i)}, y^{(i)})\}_{i=1}^N$ $f(x) = xw + b$
- Regression – Find f that minimizes our uncertainty about y given x

$$y = f(x) + n$$

- Minimizing Mean Squared Error = Minimizing Negative Log-Likelihood

$$\operatorname{argmin}_f \frac{1}{N} \sum_{i=1}^N (y^{(i)} - f(x^{(i)}))^2$$


Linear Regression

$$\operatorname{argmin}_w \frac{1}{N} \sum_{i=1}^N (y^{(i)} - w^T \hat{x}^{(i)})^2 = \operatorname{argmin}_w \frac{1}{N} \|y - Xw\|^2$$

Loss/Cost Function

Where

- $y = [y^{(1)}, \dots, y^{(N)}]^T \in \mathbb{R}^{N \times 1}$ and
- $X = [\hat{x}^{(1)}, \dots, \hat{x}^{(N)}]^T \in \mathbb{R}^{N \times (d+1)}$ (here $d = 1$)
- $w = [w_0, w_1, \dots, w_d]^T \in \mathbb{R}^{d+1}$



Linear Regression

$$J(w) = \frac{1}{N} \sum_{i=1}^N \left(y^{(i)} - w^T \hat{x}^{(i)} \right)^2$$

Compute the minimum value? How to do it in Math?

Find points where gradient = 0

Linear Regression

$$J(w) = \frac{1}{N} \sum_{i=1}^N \left(y^{(i)} - w^T \hat{x}^{(i)} \right)^2$$

$$J(W) = \frac{1}{N} (Y - XW)^T (Y - XW)$$

<https://www.math.uwaterloo.ca/~hwolkowi/matrixcookbook.pdf>

$$J(W) = \frac{1}{N} (Y^T Y - 2Y^T XW + W^T X^T XW)$$


$$\nabla J(W) = \frac{\partial}{\partial W} \left[\frac{1}{N} (Y^T Y - 2Y^T XW + W^T X^T XW) \right] \quad X^T XW = X^T Y$$

$$W^* = (X^T X)^{-1} X^T Y$$

$$\nabla J(W) = -\frac{2}{N} X^T Y + \frac{2}{N} X^T XW$$

Linear Regression

$$J(w) = \frac{1}{N} \sum_{i=1}^N \left(y^{(i)} - w^T \hat{x}^{(i)} \right)^2$$



If the gradient of a function is non-zero at a point, the direction of the gradient is the direction in which the function increases most quickly

Linear Regression

$$J(w) = \frac{1}{N} \sum_{i=1}^N \left(y^{(i)} - w^T \hat{x}^{(i)} \right)^2$$

$$\nabla J(w) = \frac{\partial J(w)}{\partial w}$$

$$\nabla J(w) = -\frac{2}{N} \sum_{i=1}^N (y^{(i)} - w^T \hat{x}^{(i)}) \hat{x}^{(i)}$$

$$J(w) = \frac{1}{N} \sum_{i=1}^N (y^{(i)} - w^T \hat{x}^{(i)})^2$$



$$w_0 := w_0 + \frac{2\alpha}{N} \sum_{i=1}^N (y^{(i)} - (w_0 + w_1 x_1^{(i)}))$$

$$\frac{\partial J(w)}{\partial w} = \frac{1}{N} \sum_{i=1}^N 2(y^{(i)} - w^T \hat{x}^{(i)}) (-\hat{x}^{(i)})$$


$$w_1 := w_1 + \frac{2\alpha}{N} \sum_{i=1}^N (y^{(i)} - (w_0 + w_1 x_1^{(i)})) x_1^{(i)}$$

$$= -\frac{2}{N} \sum_{i=1}^N (y^{(i)} - w^T \hat{x}^{(i)}) \hat{x}^{(i)}$$

Local Minimum vs Global Minimum



Linear Regression




$$f(x) = w_0x + b$$

Slope

Intercept

Problem of Linear Regression



NonLinear Regression

- So far, we have been using a linear function for regression:

$$f(x) = w^T x + w_0 = \sum_{i=0}^d w_i x_i \quad (\text{Assuming } x_0 = 1)$$

- Lets generalize this model:

$$f(x) = \sum_{i=0}^M w_i \phi_i(x) = w^T \phi(x)$$

where ϕ_i are fixed “basis” functions.

- For linear regression $M = d$, $\phi_i(x) = x_i$.

NonLinear Regression

$$f(x) = \sum_{i=0}^M w_i \phi_i(x) = w^T \phi(x)$$



E.g., Polynomial Regression:

- 1D Polynomial Regression, $\phi(x) = [1, x, x^2, x^3]$:

$$\operatorname{argmin}_w \frac{1}{N} \sum_{n=1}^N (w^T \phi(x^{(i)}) - y^{(i)})^2$$

To avoid confusion, note that: $\phi(x^{(i)}) = [1, x^{(i)}, (x^{(i)})^2, (x^{(i)})^3]$

$$f(x^{(i)}) = w_0 + w_1 x^{(i)} + w_2 (x^{(i)})^2 + w_3 (x^{(i)})^3$$



NonLinear Regression

Loss:

$$\operatorname{argmin}_w \frac{1}{N} \sum_{n=1}^N (w^T \phi(x^{(i)}) - y^{(i)})^2 = \operatorname{argmin}_w \|\Phi w - y\|^2$$

Where $\Phi = [\phi(x^{(1)}), \dots, \phi(x^{(N)})]^T \in \mathbb{R}^{N \times M}$ and $w \in \mathbb{R}^M$.

Optimization:

1. Closed form solution: $w^* = (\Phi^T \Phi)^{-1} \Phi^T y$


2. Gradient descent: $w^{(t)} = w^{(t-1)} - \epsilon \nabla_w Loss(w^{(t-1)})$

$$J(w) = \frac{1}{N} \sum_{i=1}^N (y^{(i)} - w^T \hat{x}^{(i)})^2$$

$$W^* = (X^T X)^{-1} X^T Y$$

What is the problem of
Nonlinear Regression?

NonLinear Regression



The basis function is all fixed!

Can we learn the basis function?

NonLinear Regression

- Lets first look at what the learning problem might look like:



$$\underset{w}{\operatorname{argmin}} \sum_i \left(\left(\sum_j w_j \phi_j(x^{(i)}) \right) - y^{(i)} \right)^2$$

Neural Networks do this for us!

What things are learned here?
What things are fixed here?

NonLinear Regression

1-layer Multi-layer Perceptron



$$\sum_{j=0}^{d^1} w_j^1 \phi_j(x^1)$$

NonLinear Regression

- Lets first look at what the learning problem might look like:


$$\operatorname{argmin}_{w, \{\phi_j\}_{j=1}^M} \sum_i \left(\left(\sum_j w_j \phi_j(x^{(i)}) \right) - y^{(i)} \right)^2$$

Neural Networks do this for us!

What things are learned here?
What things are fixed here?


NonLinear Regression

1-hidden layer Multi-layer Perceptron



NonLinear Regression


1-hidden layer Multi-layer Perceptron



$$f(x) = \sum_{m=1}^M w_m^1 \phi_m(x)$$

Example

- Activations at hidden layers can be viewed as features extracted as functions of inputs
- Every hidden layer represents a level of abstraction
 - *Complex features are compositions of simpler features*




- Number of layers is known as depth of ANN
 - *Deeper networks express complex hierarchy of features*

Question?



1. "Judge a man by his questions rather than by his answers."
– Voltaire
 2. "If I had an hour to solve a problem, I'd spend 55 minutes thinking about the problem and 5 minutes thinking about solutions."
– Albert Einstein
 3. "The art and science of asking questions is the source of all knowledge."
– Thomas Berger
 4. "Asking the right questions takes as much skill as giving the right answers."
– Robert Half
 5. "The wise man doesn't give the right answers, he poses the right questions."
– Claude Lévi-Strauss
 6. "Great questions make great companies."
– Peter Drucker

Linear Regression




$$\sigma(w^T x + b)$$

$$\operatorname{argmin}_w \frac{1}{N} \sum_{i=1}^N (y^{(i)} - w^T \hat{x}^{(i)})^2 = \operatorname{argmin}_w \frac{1}{N} \|y - Xw\|^2$$


Loss/Cost Function

Analytical
Solution


Gradient
Descent



$$\nabla J(W) = -\frac{2}{N} X^T Y + \frac{2}{N} X^T X W \quad W^* = (X^T X)^{-1} X^T Y$$



Linear Regression



$$y = \sigma(w^T x + b)$$

$$y = \sigma(w^T \phi(x) + b)$$

$$\nabla J(W) = -\frac{2}{N} X^T Y + \frac{2}{N} X^T X W$$

$$\nabla J(W) = -\frac{2}{N} \phi(X)^T Y + \frac{2}{N} \phi(X)^T \phi(X) W$$

$$W^* = (X^T X)^{-1} X^T Y$$

$$W^* = (\phi(X)^T \phi(X))^{-1} \phi(X)^T Y$$

Linear Regression

$$J(w) = \frac{1}{N} \sum_{i=1}^N \left(y^{(i)} - w^T \hat{x}^{(i)} \right)^2$$

$$\nabla J(w) = \frac{\partial J(w)}{\partial w}$$

$$\nabla J(w) = -\frac{2}{N} \sum_{i=1}^N (y^{(i)} - w^T \hat{x}^{(i)}) \hat{x}^{(i)}$$

$$J(w) = \frac{1}{N} \sum_{i=1}^N (y^{(i)} - w^T \hat{x}^{(i)})^2$$


$$w_0 := w_0 + \frac{2\alpha}{N} \sum_{i=1}^N (y^{(i)} - (w_0 + w_1 x_1^{(i)}))$$

$$\frac{\partial J(w)}{\partial w} = \frac{1}{N} \sum_{i=1}^N 2(y^{(i)} - w^T \hat{x}^{(i)}) (-\hat{x}^{(i)})$$

$$w_1 := w_1 + \frac{2\alpha}{N} \sum_{i=1}^N (y^{(i)} - (w_0 + w_1 x_1^{(i)})) x_1^{(i)}$$

$$= -\frac{2}{N} \sum_{i=1}^N (y^{(i)} - w^T \hat{x}^{(i)}) \hat{x}^{(i)}$$

Linear Regression




$$y = w_1 x + w_0$$

$$w_0 := w_0 + \frac{2\alpha}{N} \sum_{i=1}^N (y^{(i)} - (w_0 + w_1 x_1^{(i)}))$$

$$w_1 := w_1 + \frac{2\alpha}{N} \sum_{i=1}^N (y^{(i)} - (w_0 + w_1 x_1^{(i)})) x_1^{(i)}$$

Non-Linear Regression



$$y = \sigma(w^T x + b)$$

$$y = \sigma(w^T \phi(x) + b)$$

$$\nabla J(W) = -\frac{2}{N} X^T Y + \frac{2}{N} X^T X W$$


$$\nabla J(W) = -\frac{2}{N} \phi(X)^T Y + \frac{2}{N} \phi(X)^T \phi(X) W$$

$$W^* = (X^T X)^{-1} X^T Y$$

$$W^* = (\phi(X)^T \phi(X))^{-1} \phi(X)^T Y$$


Multi-NonLinear Regression

1-hidden layer Multi-layer Perceptron




Multi-NonLinear Regression

1-hidden layer Multi-layer Perceptron




Classification


Regression (y continuous)



Classification (y categorical/discrete)




Classification



What loss should we use?


$$\mathcal{L}(w) = - \sum_{i=1}^N y^{(i)} \log(f(x^{(i)})) + (1 - y^{(i)}) \log(1 - f(x_i))$$

MINIST Dataset



A 256x256 (RGB) image \Rightarrow ~200K dimensional input x

MLP for MNIST



A fully connected network would need a very large number of parameters, very likely to overfit the data

CNN

Class 1



:



Class 2




:




- The features of these classes are spatially local.
- Translation does not change the identity of the classes, i.e., we require a translation equivariant model.
- MLP is not a good match to this problem

Can represent a small region with fewer parameters





CNN





- The features of these classes are spatially local.
- Translation does not change the identity of the classes, i.e., we require a translation equivariant model.
- MLP is not a good match to this problem

Can represent a small region with fewer parameters




CNN



From U-Waterloo's
Deep Learning Course

What about training a lot of such “small” detectors and each detector must “move around”.


Convolution Layer



CNN

Convolution Layer


consider a second, green filter



32x32x3 image
5x5x3 filter


convolve (slide) over all
spatial locations

activation maps



CNN

For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



We stack these up to get a “new image” of size 28x28x6!

CNN

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

: :

Each filter detects a small pattern (3 x 3).

CNN

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

Dot
product




1	-1	-1
-1	1	-1
-1	-1	1


Filter 1

6 x 6 image

CNN



6 x 6 image



Convolutions (typically with *prespecified* filters) are a common operation in many computer vision applications



Original image z



Gaussian blur



$$z * \begin{bmatrix} 1 & 4 & 7 & 4 & 1 \\ 4 & 16 & 26 & 16 & 4 \\ 7 & 26 & 41 & 26 & 7 \\ 4 & 16 & 26 & 16 & 4 \\ 1 & 4 & 4 & 4 & 1 \end{bmatrix} / 273$$

$$\left(\left(z * \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix} \right)^2 + \left(z * \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix} \right)^2 \right)^{\frac{1}{2}}$$



Image gradient

CNN



Convolution is a linear operator

We need to follow convolution with a nonlinearity (e.g., ReLU) to get nonlinear functions.

CNN

bird



$$\begin{array}{|c|c|c|} \hline 1 & -1 & -1 \\ \hline -1 & 1 & -1 \\ \hline -1 & -1 & 1 \\ \hline \end{array}$$



Filter 1

$$\begin{array}{|c|c|c|} \hline -1 & 1 & -1 \\ \hline -1 & 1 & -1 \\ \hline -1 & 1 & -1 \\ \hline \end{array}$$


Filter 2




bird




CNN



Equivariant and Invariant




input image



input vector

Equivariant and Invariant




must learn shift invariance from data!




Code Demo

Equivariant and Invariant

\mathfrak{G} -invariance $f(\rho(g)x) = f(x)$



\mathfrak{G} -equivariance $f(\rho(g)x) = \rho(g)f(x)$



Equivariant and Invariant





Image Convolution is ?

Deep Neural Networks

20th Century Zoo of Neural Network Architectures



DNNs

IMAGENET




Speech data




Natural language processing (NLP)

...





Grid games



Deep neural nets that exploit:


- translation equivariance (weight sharing)
- hierarchical compositionality



DNNs


A lot of real-world data does not “live” on grids

Social networks
Citation networks
Communication networks
Multi-agent systems




**Protein interaction
networks**

From CNN to GNN



Knowledge graphs



Road maps



* slide from Thomas Kipf, **University of Amsterdam**

Sequential Data


The teacher told the student that he was brilliant.

The student told the teacher that he was brilliant.


**You read this sentence from left to right and
understand the sentence**

Sequential Data - RNN

What if we have sequences of variable lengths, like sentences or videos that we would like to analyze?




Sequential Data - RNN




Token embedding –
DNA of the Token

Sequential Data - RNN




Sequential Data - RNN




Sequential Data - RNN

What if we have sequences of variable lengths, like sentences or videos that we would like to analyze?




Time or Progression

How would you setup
the model (token
embedding, MLP layer)



- $h_t = \sigma(W^T \begin{bmatrix} x_t \\ h_{t-1} \end{bmatrix})$
- $y_t = F(h_t)$

Sequential Data - RNN



Sequential Data - RNN

Image Captioning


- Given an image, produce a sentence describing its contents
- Inputs: Image feature (from a CNN)
- Outputs: Multiple words (let's consider one sentence)



: The dog is hiding


Sequential Data - RNN

Image Captioning



: Arun Mallya

Sequential Data - RNN



$$y_1 = f_1(x; W_1)$$

$$y_2 = f_2(y_1; W_2)$$

$$C = \text{Loss}(y_2, y_{GT})$$

Find $\frac{\partial C}{\partial W_1}, \frac{\partial C}{\partial W_2}$

$$\frac{\partial C}{\partial W_2} = \left(\frac{\partial C}{\partial y_2} \right) \left(\frac{\partial y_2}{\partial W_2} \right)$$

$$\frac{\partial C}{\partial W_1} = \left(\frac{\partial C}{\partial y_1} \right) \left(\frac{\partial y_1}{\partial W_1} \right)$$

$$= \left(\frac{\partial C}{\partial y_2} \right) \left(\frac{\partial y_2}{\partial y_1} \right) \left(\frac{\partial y_1}{\partial W_1} \right)$$

Sequential Data - RNN

$$h_t = f(Wh_{t-1} + b)$$

$$\left\| \frac{\partial L}{\partial h_0} \right\| \approx \sigma_{\max}^T \left\| \frac{\partial L}{\partial h_T} \right\|$$


- Largest singular value $> 1 \rightarrow \text{Exploding gradients}$
 - Slight error in the late time steps causes drastic updates in the early time steps \rightarrow Unstable learning
- Largest singular value $< 1 \rightarrow \text{Vanishing gradients}$
 - Gradients passed to the early time steps is close to 0. \rightarrow Uninformed correction

Any other problem with RNN?

Sequential Data - RNN

Input – Output Scenarios

Single - Single



Feed-forward Network

Single - Multiple






Image Captioning

Multiple - Single



Sentiment Classification

Multiple - Multiple



Translation

Sequential Data - Transformer

The animal didn't cross the street because it was too tired .

The animal didn't cross the street because it was too wide .

Key

Value

Query

Sequential Data - Transformer

Think of YouTube.

*First you enter your **Query** in the search bar. Then your **Query** is compared against a set of **Keys** (in this case, video titles, tags and descriptions etc. within the YouTube database). After this, YouTube proceeds to retrieve the videos that best match your **Query**. These video results are referred to as **Values**.*

Key	Value	Query
------------	--------------	--------------

I would like to **taste** the local food in **France**.

Sequential Data - Transformer


I would like to taste the local food in France.

$$Q = XW_Q \quad K = XW_K \quad V = XW_V$$


$Q \cdot K = QK^T =$

	I	would	like	to	taste	the	local	food	in	France
I	<i>h</i>									
would		<i>h</i>								
like			<i>h</i>							
to				<i>h</i>						
taste					<i>h</i>		<i>h</i>			
the						<i>h</i>				
local							<i>h</i>		<i>h</i>	
food							<i>h</i>	<i>h</i>		
in									<i>h</i>	
France							<i>h</i>		<i>h</i>	

$$\text{softmax}\left(\frac{Q_i K^T}{\sqrt{d_k}}\right) = \frac{e^{\left(\frac{Q_i K^T}{\sqrt{d_k}}\right)}}{\sum_{j=1}^{d_k} e^{\left(\frac{Q_j K^T}{\sqrt{d_k}}\right)}}$$



Sequential Data - Transformer




$$\alpha_i = \frac{\exp(f_{attn}(\text{key}_i, \text{query}))}{\sum_j \exp(f_{attn}(\text{key}_j, \text{query}))}, \quad \text{out} = \sum_i \alpha_i \cdot \text{value}_i$$

Sequential Data - Transformer

The scaled dot product attention allows a network to attend over a sequence. However, often there are multiple different aspects a sequence element wants to attend to, and a single weighted average is not a good option for it. This is why we extend the attention mechanisms to multiple heads

$$\text{Multihead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$$
$$\text{where } \text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$$



$$W_{1\dots h}^Q \in \mathbb{R}^{D \times d_k}, W_{1\dots h}^K \in \mathbb{R}^{D \times d_k}, W_{1\dots h}^V \in \mathbb{R}^{D \times d_v}, \text{and } W^O \in \mathbb{R}^{h \cdot d_v \times d_{out}}$$

Sequential Data - Transformer

I would like to taste the local food in France.

$$Q = XW_Q \quad K = XW_K \quad V = XW_V$$

Self-Attention

Attention weight between the tokens corresponding to "Life" and "short"

	Life	is	short	eat	desert	first
Life	0.17	0.13	0.18	0.16	0.15	0.18
is	0.03	0.68	0.02	0.08	0.14	0.02
short	0.19	0.06	0.25	0.14	0.11	0.23
eat	0.15	0.21	0.14	0.16	0.17	0.14
desert	0.13	0.27	0.11	0.16	0.18	0.12
first	0.19	0.02	0.31	0.11	0.07	0.27


Causal-Attention

In the row for corresponding to "Life", mask out all words that come after "Life"


	Life	is	short	eat	desert	first
Life	0.17	0.13	0.18	0.16	0.15	0.18
is	0.03	0.68	0.02	0.08	0.14	0.02
short	0.19	0.06	0.25	0.14	0.11	0.23
eat	0.15	0.21	0.14	0.16	0.17	0.14
desert	0.13	0.27	0.11	0.16	0.18	0.12
first	0.19	0.02	0.31	0.11	0.07	0.27



Sequential Data - Transformer




$$\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2$$
$$x = \text{LayerNorm}(x + \text{FFN}(x))$$



$$\mu^l = \frac{1}{H} \sum_{i=1}^H a_i^l \quad \sigma^l = \sqrt{\frac{1}{H} \sum_{i=1}^H (a_i^l - \mu^l)^2}$$

Sequential Data - Transformer

Cross-attention



Sequential Data - Transformer

Multi-Attention

Self-Attention

Sentiment Analysis

Causal-Attention

Next-token
Prediction

Cross-Attention

Translation

Sequential Data - Transformer

The teacher told the student that he was brilliant.

The student told the teacher that he was brilliant.

**Two sentences using exactly the same tokens but end
up with very different meaning**

Sequential Data - Transformer

0 :	0	0	0	0	8 :	1	0	0	0
1 :	0	0	0	1	9 :	1	0	0	1
2 :	0	0	1	0	10 :	1	0	1	0
3 :	0	0	1	1	11 :	1	0	1	1
4 :	0	1	0	0	12 :	1	1	0	0
5 :	0	1	0	1	13 :	1	1	0	1
6 :	0	1	1	0	14 :	1	1	1	0
7 :	0	1	1	1	15 :	1	1	1	1

Sequential Data - Transformer

$$P(k, 2i) = \sin\left(\frac{k}{n^{2i/d}}\right)$$


k: Position of an object in the input sequence,

$$P(k, 2i + 1) = \cos\left(\frac{k}{n^{2i/d}}\right)$$


d: Dimension of the output embedding space

Sequence	Index of token, k	Positional Encoding Matrix with $d=4$, $n=100$			
		$i=0$	$i=0$	$i=1$	$i=1$
I	0	$P_{00}=\sin(0) = 0$	$P_{01}=\cos(0) = 1$	$P_{02}=\sin(0) = 0$	$P_{03}=\cos(0) = 1$
am	1	$P_{10}=\sin(1/1) = 0.84$	$P_{11}=\cos(1/1) = 0.54$	$P_{12}=\sin(1/10) = 0.10$	$P_{13}=\cos(1/10) = 1.0$
a	2	$P_{20}=\sin(2/1) = 0.91$	$P_{21}=\cos(2/1) = -0.42$	$P_{22}=\sin(2/10) = 0.20$	$P_{23}=\cos(2/10) = 0.98$
Robot	3	$P_{30}=\sin(3/1) = 0.14$	$P_{31}=\cos(3/1) = -0.99$	$P_{32}=\sin(3/10) = 0.30$	$P_{33}=\cos(3/10) = 0.96$


Sequential Data – Vision Transformer





Sequential Data – GPT Transformer




Sequential Data - Transformer




Sequential Data - Transformer




Sequential Data - Transformer



Sequential Data – Transformer - Training




Sequential Data – Transformer - Inference



**Inference Procedure
(KV cache)**


Sequential Data – Transformer - Inference



Autoregressive Nature

$$P(\hat{\mathbf{y}}_i | \mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m, \hat{\mathbf{y}}_1, \hat{\mathbf{y}}_2, \dots, \hat{\mathbf{y}}_{i-1})$$

Sequential Data – Transformer - Inference




Sequential Data – Transformer - Inference

$$\hat{\mathbf{y}}_i = \arg \max_{\hat{\mathbf{y}}} P(\hat{\mathbf{y}} | \mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m, \hat{\mathbf{y}}_1, \hat{\mathbf{y}}_2, \dots, \hat{\mathbf{y}}_{i-1})$$



Greedy decoding is computationally efficient and easy to implement.
It does not explore alternative paths that might lead to more globally optimal sequences.

Sequential Data – Transformer - Inference



<https://medium.com/@lmpo/mastering-langs-a-guide-to-decoding-algorithms-c90a48fd167b>

Deep Neural Networks

20th Century Zoo of Neural Network Architectures

